

ACCOUNT-BASED MARKETING

AI Agents for ABM.

The five agents and one workflow that make Tier 1 account orchestration possible at scale.

WRITTEN BY

Erik R. Miller

SENIOR B2B MARKETING EXECUTIVE · FRACTIONAL CMO

5

AGENTS

1

WORKFLOW

23

PAGES

AI AGENTS FOR ABM
THE OPERATOR PLAYBOOK · VOL. 04

FIRST EDITION, 2026
© 2026 ERIK R. MILLER · E.R.M. ADVISORY
NEW YORK CITY

ERIKRMILLER.COM
THE OPERATOR NEWSLETTER · THEOPERATOR-ERM.BEEHIV.COM

This Field Note is one volume in the Operator Playbook series.
Adapted from the flagship essay AI Agents for Marketing Teams, originally published on erikrmiller.com.
Permission granted to forward, print, and share this document in full. Excerpts welcomed with attribution.

ABM is the function AI agents were built for.

Account-based marketing is where the gap between what humans can run manually and what a marketing org needs to deliver shows up first. Tier 1 motion across thirty accounts is a stretch for a senior ABMer. Across fifty accounts it is a fantasy. Across two hundred accounts it is a slide in a deck that no one is executing.

AI agents are the way ABM scales without exploding headcount. Account selection refreshes weekly off intent and fit grades. Orchestration sequences run across every Tier 1 in parallel. Competitor intelligence flows into the messaging in real time. Sales enablement assets generate themselves at the per-account level.

This playbook covers the five agents that anchor the ABM stack, plus the Tier 1 motion workflow that ties them together. If you run ABM, this is the document to anchor your 2026 plan.

WHO THIS IS FOR

For ABM leads, enterprise marketing directors, and the operators running Tier 1 and Tier 2 account programs.

CONTENTS

Inside this *Field Note*.

5 agents and 1 hero workflow.

THE AGENT ROSTER

AGENT 01	The Account Selection Agent
AGENT 02	The Intent Signal Agent
AGENT 03	The ABM Orchestration Agent
AGENT 04	The Competitor Intelligence Agent
AGENT 05	The Sales Enablement Agent

HERO WORKFLOW

WORKFLOW	Tier 1 ABM motion design
-----------------	--------------------------------

OPERATIONS & EXECUTION

—	Pre-flight checklist
—	The 30-60-90 Roadmap
—	Pitfalls & remediation

WHERE THIS FITS

—	Where this fits.
---	-----------------------

The ABM. Agents.

AGENT 01

The Account Selection Agent

PURPOSE. Builds and refreshes target account lists by tier, drawing on the ICP Research Agent's grades and the Intent Signal Agent's signal stack. Outputs the weekly Tier 1, Tier 2, and Tier 3 lists that drive the ABM motion. Most ABM programs fail because account selection is wrong; [this is where that failure starts](#).

INPUTS. ICP grades; first-party intent data; third-party intent feeds; funding, leadership, and hiring news; public-record signals; CRM activity from the last 90 days; current sales capacity by tier.

OUTPUTS. Tier-segmented account lists; weekly tier-movement report (accounts entering/leaving each tier with reasoning); a "watchlist" of accounts approaching tier-promotion thresholds.

MEMORY. Account-level history (when entered each tier, motion run, outcome); tier-promotion accuracy by signal type; recent disqualification reasoning.

HUMAN OVERSIGHT. ABM lead and sales leadership review the Tier 1 list weekly. Sales has veto power on Tier 1 entry to keep the list operationally honest.

COMMON MISTAKES. Over-rotating Tier 1 accounts week-to-week, exhausting sales attention. Letting tier-promotion be triggered by single-source intent (third-party intent alone is the noisiest predictor). Failing to cap Tier 1 at a number sales can actually run motion against.

HOW TO DEPLOY

The Account Selection Agent operationalizes the ICP into tier-segmented account lists. It runs weekly and produces the only account list marketing and sales should be working off.

- 01** Define the tier model before deploying. Tier 1, 2, 3, 4 with explicit criteria for each. Tier 1 typically means concentrated motion, named ABMer, custom content. The agent does not get to invent the tier definitions.
- 02** Connect the agent to the ICP Research Agent's grading output, first-party intent data sources, and at least one third-party intent feed. Without all three, the agent will over-rotate on whichever signal is available.
- 03** Produce the first list in audit mode. Compare against the list sales is currently working. Investigate every account where the agent and sales disagree. The disagreements are where the value is.
- 04** Hand the agent's Tier 1 list to the ABM lead and Sales VP for joint weekly review. Sales has veto power on Tier 1 promotions. This is not a policy debate, it is a deployment requirement.
- 05** Promote to live state once the joint review has produced two consecutive weeks of clean lists. Long-run the weekly tier-movement report as the standing artifact.

SUCCESS METRICS · WHAT TO MEASURE

Tier 1 win rate	Percent of accounts the agent promoted to Tier 1 that closed-won within the typical sales cycle. Target: at least 2x average company win rate.
Tier accuracy at promotion time	Tier-promotion accuracy measured against eventual outcome. Target: 70%+ for Tier 1, 50%+ for Tier 2.
Sales adoption	Percent of sales-pursued accounts that appear on the agent's current tier list. Target: > 85%.
Veto rate	Percent of agent-promoted Tier 1 accounts vetoed by sales each week. Target: < 15%. A higher rate means the ICP or the tier criteria are off.

The Intent Signal Analysis Agent

PURPOSE. Synthesizes first-party and third-party intent signals into account-level "what changed and what to do" briefs. Distinguishes high-signal events from noise. [Most intent data is mostly hype](#); this agent's job is to enforce that judgment programmatically.

INPUTS. First-party intent (web visits, content engagement, demo requests, pricing-page hits); third-party intent feeds; news and funding feeds; LinkedIn engagement; community signals; product-usage signals where applicable.

OUTPUTS. Daily account-level signal briefs (top 20 accounts with new signals); signal-to-action recommendations (what motion fits each signal type); a weekly "signal noise" report flagging which third-party feeds added value vs. just volume.

MEMORY. Signal-to-outcome correlation history (which signals predicted pipeline, which did not); account-level signal history; signal-decay calibration (how quickly each signal type loses predictive value).

HUMAN OVERSIGHT. ABM lead reviews top-account signals daily. Quarterly review of signal-source ROI — sources that fail to predict pipeline are dropped.

COMMON MISTAKES. Treating volume of signals as a signal. Failing to calibrate signal decay. Buying a third-party intent feed and never auditing whether it added pipeline.

HOW TO DEPLOY

The Intent Signal Analysis Agent converts noisy intent feeds into prioritized action. Without this agent, intent data is a distraction. With it, intent data becomes the highest-leverage input to the Account Selection Agent and the ABM Orchestration Agent.

- 01** Define what counts as a meaningful intent signal for your business. Generic content topic intent rarely qualifies. Buyer-side comparison searches, RFP-language queries, and review-site visits qualify. The agent operates only on signals you have defined as meaningful.
- 02** Connect the agent to your first-party signals (gated content, demos, ad clicks), at least one third-party intent feed, and the Account Selection Agent's tier list. Intent signals on non-ICP accounts get filtered out, not investigated.
- 03** Run the agent in advisory mode for three weeks. It produces a weekly Top 20 signal report with recommended action per account. ABM lead reviews and decides what is acted on. The agent does not yet trigger workflows.
- 04** After three weeks, promote the agent to trigger ABM Orchestration motion on the highest-confidence Tier 1 signals. Lower-confidence signals stay in advisory mode for further development.
- 05** Tune the signal-to-action thresholds monthly. Every month the agent reports false-positive rate, missed-signal rate, and time-from-signal-to-pipeline. Tune the thresholds against this report.

SUCCESS METRICS • WHAT TO MEASURE

Signal-to-pipeline conversion

Percent of agent-flagged signals that produce a meeting or pipeline entry within 30 days. Target: > 12%.

False positive rate

Percent of high-confidence signals where ABM motion produced no engagement. Target: < 30%.

Time-to-action

Median time from signal detection to first outbound touch. Target: < 48 hours for Tier 1.

Coverage

Percent of pipeline-creating accounts where the agent caught a signal in the 60 days prior to opportunity creation. Target: > 50%.

The ABM Orchestration Agent

PURPOSE. Coordinates multi-channel ABM motions for Tier 1 and Tier 2 accounts: content touch points, paid programs, sales outreach, executive engagement. Sequences touches so each account experiences a coherent message, not scattered noise. The single most operationally complex agent in the system, and the most leveraged for ABM-led organizations.

INPUTS. Tier 1 and Tier 2 account lists; account-specific situational signals; persona maps; channel-cost models; current campaign assets; sales cadence and capacity; [the ABM tier model](#).

OUTPUTS. 30/60/90-day motion plan per account; channel-touch-sequence map; asset list (custom and standard); sales talk track tied to account-specific context; weekly motion-fidelity report.

MEMORY. Per-account motion history; response patterns; channel performance per account-tier; prior orchestration mistakes.

HUMAN OVERSIGHT. ABM lead approves Tier 1 motion plans before activation. Weekly review of execution fidelity (planned vs. actual touches). Monthly ABM-pipeline review.

COMMON MISTAKES. Generating touch sequences that look comprehensive but cost more than the deal is worth. Misaligning with sales cadence. Treating Tier 2 like Tier 1 (different motions, different economics).

HOW TO DEPLOY

The ABM Orchestration Agent is the operating system of account-based marketing. It runs Tier 1 motion at a scale no human ABMer can match while keeping the per-account discipline that ABM requires. Deploy it after the Account Selection Agent has stabilized.

- 01** Build the per-tier playbook. Tier 1 means a defined sequence of touches (custom content, executive intro, named-account ad, sales outreach), with defined timing and defined assets. The agent runs this playbook, it does not invent one.
- 02** Connect the agent to the Account Selection Agent's tier list, the Intent Signal Analysis Agent's triggers, the Content Strategy Agent's asset library, and the Sales Enablement Agent's per-account materials.
- 03** Run the first 30 days against a held-out sample. Pick fifteen Tier 1 accounts. The agent orchestrates the motion. The ABM lead audits every touch before it ships for the first two weeks.
- 04** Promote to full Tier 1 coverage in week 5. The audit-every-touch discipline drops to audit-every-motion. ABM lead reviews the agent's motion plans, not each individual touch.
- 05** Quarterly, expand the agent to Tier 2 once Tier 1 motion fidelity holds above 90%. Tier 3 and Tier 4 stay in advisory mode unless explicit decision to scale.

SUCCESS METRICS • WHAT TO MEASURE

Tier 1 motion fidelity	Percent of planned touches that execute on schedule and against the playbook. Target: > 90%.
Tier 1 engaged-account rate	Percent of Tier 1 accounts producing a measurable engagement (meeting, content consumption, intent spike) within 60 days. Target: > 25%.
ABM-sourced pipeline	Dollars of pipeline sourced from Tier 1 + Tier 2 accounts per quarter. Target: trend up, with a measurable baseline by quarter 1.
Sales co-ownership	Percent of Tier 1 accounts with a named AE actively participating in the motion. Target: 100%.

The Competitor Intelligence Agent

PURPOSE. Tracks competitor positioning, product, pricing, content, and field motion. Produces battlecards, threat alerts, and quarterly competitive briefs. The agent's most important rule is to never invent. Hallucinated competitor claims are the fastest way to destroy sales trust in the system.

INPUTS. Competitor websites, public earnings, news, hiring signals, review-site verbatims, win/loss interviews, sales-team field notes, social media positioning shifts.

OUTPUTS. Per-competitor positioning briefs; quarterly battlecard refresh; threat alerts (new product launches, pricing moves, leadership changes); message-overlap analysis showing where the brand and a competitor are saying the same thing.

MEMORY. Competitor positioning over time; verified vs. unverified claims (with sources); historical accuracy of threat alerts.

HUMAN OVERSIGHT. Sales enablement reviews battlecards before distribution. Any competitor claim must include a verifiable source citation. Unsourced claims are flagged, never published.

COMMON MISTAKES. Hallucinating competitor features, pricing, or customer wins. Treating press coverage as truth. Over-indexing on one loud competitor and missing the quiet one taking share.

HOW TO DEPLOY

The Competitor Intelligence Agent has one rule above all others. It does not invent. Hallucinated competitor claims destroy sales trust faster than any other agent failure. Deploy this agent with the most cautious oversight model in the stack.

- 01** Define the competitor set explicitly. Top three direct, two adjacent, one emerging. The agent works only against this defined set. Anything outside the set requires a deliberate scope change.
- 02** Connect the agent to public sources only at first. Competitor websites, public earnings, news, hiring signals, review-site verbatims. Do not feed the agent unsourced field rumor as input.
- 03** Establish the never-invent rule in the agent's system prompt. If a fact cannot be cited to a source, the agent must either flag missing data or refuse to write. Test the rule before promoting the agent to production.
- 04** Produce the first quarterly battlecard refresh in shadow mode. The Head of Strategy compares the agent's output against the existing sales-trusted positioning. Document every divergence.
- 05** Promote to live state. The agent now produces weekly threat alerts and quarterly battlecard refreshes. The Head of Strategy approves any quarterly narrative shifts before publication.

SUCCESS METRICS · WHAT TO MEASURE

Citation rate	Percent of competitor claims in agent output that include a verifiable source link. Target: 100%.
Hallucination incidents	Number of unsourced or fabricated competitor claims found in any agent output. Target: 0. This is a kill criterion.
Threat-alert lead time	Median time from a competitor move (launch, pricing, leadership change) being publicly visible to the agent flagging it. Target: < 5 business days.
Sales adoption	Percent of active sales reps who reference the current battlecard in their last five calls. Target: > 70%.

The Sales Enablement Agent

PURPOSE. Translates marketing campaigns, ICP shifts, and competitive intelligence into sales-ready talk tracks, battlecards, sequence templates, and objection-handling scripts. [Builds the assets sales actually uses](#), not the assets that sit unread in a content portal.

INPUTS. Campaign briefs, ICP and persona docs, competitive intelligence, recent objections from sales calls, product positioning, current quarterly priorities, win/loss reasons.

OUTPUTS. SDR talk tracks, AE call openers, battlecards, sequence templates, objection-handling pages, monthly "what changed" digest for the sales team.

MEMORY. Sales-feedback log on which assets got used, which got ignored; objection patterns over time; win/loss themes.

HUMAN OVERSIGHT. Sales leadership co-owns this agent with marketing. Weekly sync on what is and is not landing.

COMMON MISTAKES. Building enablement assets sales never uses (volume without adoption). Disconnect from sales-call reality (the agent reads marketing's positioning rather than the field's reality). Failing to update when ICP shifts.

HOW TO DEPLOY

The Sales Enablement Agent is the marketing-to-revenue interface. It produces account-specific, situation-specific assets at a velocity the sales team can actually use. Deploy it after the ICP, persona, and content agents are live, because enablement built on generic foundations stays generic.

- 01** Audit the existing enablement library. Most of it will be generic decks that no AE has used in twelve months. Archive the generic. Identify the three to five asset types AEs actually use (battlecard, account-specific deck, ROI calculator, executive briefing template, RFP language).
- 02** Connect the agent to the Account Selection Agent's tier data, the Buyer Persona Agent's briefs, the Competitor Intelligence Agent's battlecards, and the CRM's account context.
- 03** Run the agent on the top 30 active Tier 1 accounts. Produce per-account assets. Have the named AE rate each asset on relevance, accuracy, and usability. Iterate until the ratings stabilize.
- 04** Promote to live state. The agent produces enablement assets on AE request inside the CRM. SLA is 24 hours or less for standard request types.
- 05** Long-run on a monthly utilization audit. Assets that go unused get retired. Assets that produce high meeting-conversion rates get reinforced as templates.

SUCCESS METRICS · WHAT TO MEASURE

AE-request fulfillment time	Median time from AE request to delivered asset. Target: < 24 hours for standard, < 72 hours for custom.
Asset utilization	Percent of agent-produced assets that get sent to a prospect within 7 days. Target: > 65%.
Meeting-conversion lift	Conversion from first meeting to second meeting on accounts using agent-produced assets vs. baseline. Target: 1.3x baseline.
AE satisfaction	Quarterly AE rating of enablement quality. Target: > 4 out of 5.

The Hero *Workflow.*

HERO WORKFLOW

Tier 1 ABM *motion design.*

How an agent stack runs Tier 1 motion across the accounts that matter.

Eight agents and two human checkpoints turn a Tier 1 account list into a coordinated motion that executes across content, ads, outbound, sales, and intent monitoring. The ABM lead edits motion plans rather than building them. The agent system runs the orchestration. The human runs the strategy.

Eight agents, two human checkpoints, one orchestrated Tier 1 motion. The ABM lead's job becomes editing motion plans, not building them. The agent system runs the orchestration so the human can run the strategy.

AGENTS INVOLVED

- Account Selection Agent (maintains the Tier 1 list)
- Intent Signal Analysis Agent (triggers motion on Tier 1 signals)
- ABM Orchestration Agent (executes the per-account playbook)
- Competitor Intelligence Agent (informs account-specific messaging)
- Content Strategy Agent (provides asset library)
- Sales Enablement Agent (produces per-account materials)
- Email Nurture Sequence Agent (runs the long-tail track)
- Analytics & Reporting Agent (tracks engagement and pipeline)

HUMAN CHECKPOINTS

- *Joint weekly Tier 1 review with Sales VP (account additions, removals, promotions)*
- *Per-motion plan review before the motion goes live*

THE WORKFLOW, STEP BY STEP

- 01** List maintenance. The Account Selection Agent updates the Tier 1 list weekly off ICP grades, intent signals, and account-level changes (leadership moves, funding rounds, technology adoption). The ABM lead and Sales VP run a joint review every Monday.
- 02** Trigger detection. The Intent Signal Analysis Agent monitors high-confidence buying signals on Tier 1 accounts and notifies the ABM Orchestration Agent.
- 03** Motion plan generation. The ABM Orchestration Agent assembles the per-account motion plan: content sequence, ad targeting, sales touchpoint cadence, escalation triggers. The plan is account-specific, not template-rendered.
- 04** Asset production. The Sales Enablement Agent produces account-specific materials. The Content Strategy Agent flags any cornerstone or supporting content relevant to the account context. The Competitor Intelligence Agent provides current-state competitive positioning.
- 05** Motion execution. The ABM Orchestration Agent runs the motion across channels. Outbound, ads, content delivery, sales nudges. Every touch is logged to the account record.
- 06** Engagement tracking. The Analytics & Reporting Agent measures account-level engagement, scores intent escalation, and flags accounts that are heating up or going dark.
- 07** Sales handoff. When an account hits engagement thresholds, the named AE is notified with a current-state summary, brief, and recommended next action.

SUCCESS METRICS

- Tier 1 motion fidelity. Target: > 90% of planned touches execute on schedule.
- Time from signal to first touch. Target: < 48 hours.
- Tier 1 engagement rate. Target: > 25% of accounts produce a measurable engagement within 60 days.
- Tier 1 pipeline creation. Target: trend up, with a defined baseline by quarter 1.
- Sales co-ownership. Target: 100% of Tier 1 accounts have a named AE actively participating.

COMMON FAILURE MODES

- Marketing-side ABM that sales does not know about. Fix: the joint Monday review is non-negotiable. Without it, the program fails inside one quarter.
- Template motion across all Tier 1 accounts. Fix: the ABM Orchestration Agent must produce per-account plans, not templated runs. If every account gets the same touches, this is not ABM, it is mass marketing.
- Engagement signal hoarding. Fix: when intent escalates, sales is notified the same day. Marketing-side engagement reporting that never reaches the AE is a process failure.

Pre-flight *checklist.*

ABM is the function where agent stacks earn their keep most visibly. It is also the function where bad agent deployment becomes mass email at a higher cost per touch. Resolve every item below before activating.



The ICP is defined across four dimensions. ABM run against a firmographic-only ICP is mass-marketing in disguise.



The tier model is explicit (Tier 1, 2, 3, 4) with named criteria per tier. Without the tier model, the Account Selection Agent has no rules to operate against.



Per-tier playbooks are documented. Tier 1 means concentrated motion, named ABMer, custom content. Tier 4 means automated and lightweight. The agents run the playbooks. The playbooks must exist on paper first.



Sales co-ownership is structured. Every Tier 1 account has a named AE who actively participates in the motion. ABM without sales adoption is broadcast email.



Intent data sources are connected (first-party plus at least one third-party feed). The Intent Signal Analysis Agent depends on these inputs.



CRM account-level data quality has been audited and remediated. Duplicate accounts, missing fields, and contact-decay are particularly damaging to ABM agents because they operate at the account level.



The marketing-influenced pipeline and ABM-sourced pipeline metrics are defined with clear methodology. ABM measurement disagreements between marketing and sales are the single most common reason ABM programs collapse.

The 30-60-90 *Roadmap.*

The ABM-stack rollout sequence. Each phase enables the next. Order matters more here than in any other Field Note because every ABM agent depends on the one before it.

DAYS 1-30

Foundation phase. Define the tier model. Document the per-tier playbook in detail (sequence of touches, asset types, sales involvement, escalation criteria). Audit the CRM and resolve account-level data quality. Connect the intent feeds. Define and align with sales leadership on the ABM measurement framework. No ABM agent ships this phase. The Foundation is the substrate.

DAYS 31-60

Deploy the Account Selection Agent in audit mode. Connect to the ICP grades and intent feeds. The agent produces the weekly tier list. Sales leadership reviews jointly with the ABM lead. Investigate every divergence between the agent's list and the list sales is currently working. Two weeks of clean joint reviews promotes the agent to live state. Begin building the per-account asset library that the ABM Orchestration Agent will draw from.

DAYS 61-90

Deploy the Intent Signal Analysis Agent and the ABM Orchestration Agent on a held-out sample. Pick fifteen Tier 1 accounts. The ABM Orchestration Agent runs the playbook end-to-end. The ABM lead audits every touch in the first two weeks, then transitions to per-motion audits. Deploy the Competitor Intelligence Agent and the Sales Enablement Agent as supporting agents.

DAY 91+

Expand to full Tier 1 coverage once motion fidelity holds above 90%. Promote the Intent Signal Analysis Agent to trigger ABM Orchestration motion on high-confidence Tier 1 signals automatically. Expand to Tier 2 in quarter 2 if Tier 1 metrics hold. Long-run the governance discipline.

Pitfalls & *remediation.*

PITFALL 01

Treating ABM as a list problem.

The Account Selection Agent produces lists. The lists are necessary, not sufficient. ABM is a motion problem. The ABM Orchestration Agent runs the motion. Teams that deploy the Selection agent but skip the Orchestration agent generate accurate lists that nobody acts on.

PITFALL 02

Marketing and sales running parallel ABM with different lists.

When the marketing-side ABM tier list and the sales-side target-account list are different, the Tier 1 motion does not work. The lists must converge. The Account Selection Agent should be the single source of truth, with joint marketing-sales review weekly.

PITFALL 03

Measurement disputes that kill the program.

The single most common ABM failure mode is not a marketing failure, it is a measurement failure. Marketing claims influence. Sales claims sourcing. The argument consumes the program. Resolve the measurement framework before deploying ABM agents.

WHERE THIS FITS

Where this fits.

The five agents and one workflow in this playbook are the ABM slice of a larger marketing operating system. ABM does not run in isolation. It draws from the strategy layer (ICP, personas), the content layer (per-tier messaging), the measurement layer (account-level attribution), and the revenue layer (sales handoff).

For the complete operating system, read the flagship: **AI Agents for Marketing Teams**. Free, no email gate.

— Erik R. Miller

Erik R. Miller

SENIOR B2B MARKETING EXECUTIVE • FRACTIONAL CMO

Web: erikrmiller.com

Newsletter: The Operator, theoperator-erm.beehiiv.com

LinkedIn: linkedin.com/in/erikrm

Facebook: facebook.com/ERMAdvisory

Instagram: instagram.com/ermadvisory

Advisory: erikrmiller.com/#contact

Tier 1 motion across *two hundred accounts* is not a roadmap. It is an agent stack.

This Field Note is one volume in the Operator Playbook series. The full system has 25 agents, four hero workflows, a governance model, and a 90-day rollout.

Six other volumes cover Demand, Content, ABM, Operations, ICP, and Sales Enablement. The flagship pulls all of it together.

NEXT STEP

Get the *full playbook*.

FLAGSHIP	Read the full 65-page playbook
SERIES	erikrmiller.com/resources
NEWSLETTER	theoperator-erm.beehiiv.com
FACEBOOK	facebook.com/ERMAdvisory
INSTAGRAM	instagram.com/ermadvisory
ADVISE	erikrmiller.com/#contact