

DEMAND GENERATION

AI Agents for *Demand* *Generation.*

*The agent stack for the people who own the
pipeline number.*

WRITTEN BY

Erik R. Miller

SENIOR B2B MARKETING EXECUTIVE · FRACTIONAL CMO

4

AGENTS

1

WORKFLOW

21

PAGES

AI AGENTS FOR DEMAND GENERATION
THE OPERATOR PLAYBOOK · VOL. 02

FIRST EDITION, 2026

© 2026 ERIK R. MILLER · E.R.M. ADVISORY

NEW YORK CITY

ERIKRMILLER.COM

THE OPERATOR NEWSLETTER · THEOPERATOR-ERM.BEEHIV.COM

This Field Note is one volume in the Operator Playbook series.

Adapted from the flagship essay AI Agents for Marketing Teams, originally published on erikrmiller.com.

Permission granted to forward, print, and share this document in full. Excerpts welcomed with attribution.

The pipeline number is the proving ground.

Demand generation is where AI in marketing earns its keep or loses the budget. The number is visible. The expectations are quarterly. The CFO is paying attention.

Three years of "AI features" gave most demand teams a worse version of what they already had. A campaign tool that wrote subject lines. A nurture platform that suggested send times. None of it knew anything about your ICP, your last quarter, or the accounts that closed.

An agent stack changes the shape of the function. Campaigns optimize themselves against real outcomes. ABM motions run across hundreds of accounts in parallel without breaking. Nurture sequences adapt to behavior the system actually understands. The demand team stops being the bottleneck for the volume the org needs.

This playbook covers four agents that anchor the demand stack, plus the Tier 1 ABM workflow that ties them together. The full system has twenty-five agents. These are the ones a demand leader builds first.

WHO THIS IS FOR

For VPs and Directors of Demand Generation, demand marketing managers, and the operators who own the marketing-sourced pipeline number.

CONTENTS

Inside this *Field Note*.

4 agents and 1 hero workflow.

THE AGENT ROSTER

AGENT 01The Campaign Optimization Agent

AGENT 02The ABM Orchestration Agent

AGENT 03The Email Nurture Agent

AGENT 04The Intent Signal Agent

HERO WORKFLOW

WORKFLOWTier 1 ABM motion design

OPERATIONS & EXECUTION

– Pre-flight checklist

– The 30-60-90 Roadmap

– Pitfalls & remediation

WHERE THIS FITS

– Where this fits.

The *Demand Generation*. Agents.

AGENT 01

The Campaign Optimization Agent

PURPOSE. Runs continuous optimization across paid, organic, and lifecycle campaigns. Identifies budget reallocation opportunities, creative fatigue, audience overlap, and channel-mix shifts. Operates inside guardrails set by the demand-gen lead.

INPUTS. Channel-level performance data; CRM pipeline contribution by source; creative library; audience definitions; budget caps; brand-safety rules; the campaign objective.

OUTPUTS. Weekly optimization brief (budget-reallocation recommendations with reasoning); creative-fatigue alerts; audience-overlap warnings; monthly channel-mix recommendation.

MEMORY. Optimization actions taken vs. outcomes; channel-mix history; what worked in similar campaigns; current quarterly objectives.

HUMAN OVERSIGHT. Demand-gen lead approves any reallocation greater than the agent's budget cap. Brand-safety constraints are non-negotiable; agent flags conflicts but never overrides.

COMMON MISTAKES. Optimizing toward conversions that are not pipeline-qualified. Killing creative before signal stabilizes. Confusing channel attribution with channel impact.

HOW TO DEPLOY

The Campaign Optimization Agent is where demand generation either scales or stops. It manages continuous optimization across paid, organic, and lifecycle motions. Deploy it last among demand agents, because it depends on every other demand agent being in place.

- 01** Define the optimization horizon. The agent is not allowed to optimize for impressions or clicks. The optimization target is qualified pipeline created or revenue influenced, attributed back through the Analytics & Reporting Agent's model.
- 02** Connect the agent to every campaign-running platform (ads, marketing automation, social schedulers), the Analytics & Reporting Agent's pipeline data, and the Ad Copy Agent's variants.
- 03** Run the agent in advisory mode for three weeks. Daily it recommends budget shifts, channel reallocations, and pacing changes. The demand lead reviews and approves manually.
- 04** Promote to bounded autonomy. The agent can move up to a defined daily budget threshold without approval. Anything above the threshold requires human sign-off. Audit the autonomous moves weekly.
- 05** Tune the bounds quarterly. As trust accumulates, the autonomy threshold rises. As mistakes happen, the threshold drops. The threshold is a real-time trust signal, not a permanent setting.

SUCCESS METRICS · WHAT TO MEASURE

Marketing-sourced pipeline efficiency	Cost per dollar of marketing-sourced pipeline vs. baseline. Target: 20% improvement within 90 days.
Channel mix accuracy	Percent of monthly budget shifts that produce a same-month or next-month pipeline lift. Target: > 60%.
Pacing discipline	Percent of campaigns hitting their monthly spend target within +/-10%. Target: > 90%.
Autonomy bound integrity	Number of unauthorized over-threshold spend events per quarter. Target: 0.

The ABM Orchestration Agent

PURPOSE. Coordinates multi-channel ABM motions for Tier 1 and Tier 2 accounts: content touch points, paid programs, sales outreach, executive engagement. Sequences touches so each account experiences a coherent message, not scattered noise. The single most operationally complex agent in the system, and the most leveraged for ABM-led organizations.

INPUTS. Tier 1 and Tier 2 account lists; account-specific situational signals; persona maps; channel-cost models; current campaign assets; sales cadence and capacity; [the ABM tier model](#).

OUTPUTS. 30/60/90-day motion plan per account; channel-touch-sequence map; asset list (custom and standard); sales talk track tied to account-specific context; weekly motion-fidelity report.

MEMORY. Per-account motion history; response patterns; channel performance per account-tier; prior orchestration mistakes.

HUMAN OVERSIGHT. ABM lead approves Tier 1 motion plans before activation. Weekly review of execution fidelity (planned vs. actual touches). Monthly ABM-pipeline review.

COMMON MISTAKES. Generating touch sequences that look comprehensive but cost more than the deal is worth. Misaligning with sales cadence. Treating Tier 2 like Tier 1 (different motions, different economics).

HOW TO DEPLOY

The ABM Orchestration Agent is the operating system of account-based marketing. It runs Tier 1 motion at a scale no human ABMer can match while keeping the per-account discipline that ABM requires. Deploy it after the Account Selection Agent has stabilized.

- 01** Build the per-tier playbook. Tier 1 means a defined sequence of touches (custom content, executive intro, named-account ad, sales outreach), with defined timing and defined assets. The agent runs this playbook, it does not invent one.
- 02** Connect the agent to the Account Selection Agent's tier list, the Intent Signal Analysis Agent's triggers, the Content Strategy Agent's asset library, and the Sales Enablement Agent's per-account materials.
- 03** Run the first 30 days against a held-out sample. Pick fifteen Tier 1 accounts. The agent orchestrates the motion. The ABM lead audits every touch before it ships for the first two weeks.
- 04** Promote to full Tier 1 coverage in week 5. The audit-every-touch discipline drops to audit-every-motion. ABM lead reviews the agent's motion plans, not each individual touch.
- 05** Quarterly, expand the agent to Tier 2 once Tier 1 motion fidelity holds above 90%. Tier 3 and Tier 4 stay in advisory mode unless explicit decision to scale.

SUCCESS METRICS · WHAT TO MEASURE

Tier 1 motion fidelity	Percent of planned touches that execute on schedule and against the playbook. Target: > 90%.
Tier 1 engaged-account rate	Percent of Tier 1 accounts producing a measurable engagement (meeting, content consumption, intent spike) within 60 days. Target: > 25%.
ABM-sourced pipeline	Dollars of pipeline sourced from Tier 1 + Tier 2 accounts per quarter. Target: trend up, with a measurable baseline by quarter 1.
Sales co-ownership	Percent of Tier 1 accounts with a named AE actively participating in the motion. Target: 100%.

The Email Nurture Sequence Agent

PURPOSE. Designs and drafts multi-step email nurture sequences (post-content download, post-event, post-demo, win-back, vertical) with branching logic and exit conditions. Owns the email layer of the buyer's journey from first touch through MQL — and replaces it with pipeline-qualified handoff because [the MQL is dead](#).

INPUTS. Persona, ICP segment, entry trigger, exit goal, prior sequence performance, brand voice samples, email engagement benchmarks.

OUTPUTS. Full sequence (subject lines, preview text, body copy, CTAs, branching logic, exit conditions, A/B test recommendations), per-email purpose statement, sequence-level performance KPIs.

MEMORY. Past sequences with engagement and conversion data, subject-line patterns that landed, exit reasons that correlated with disengagement.

HUMAN OVERSIGHT. Lifecycle lead approves new sequences. Brand Voice Governance Agent scores every email before activation. Quarterly nurture audit.

COMMON MISTAKES. Sequence length disconnected from sales cycle. Subject lines optimized for opens at the cost of trust. CTA escalation pace too aggressive for the persona's buying stage.

HOW TO DEPLOY

The Email Nurture Sequence Agent runs the long-tail conversion motion. It is not a one-shot copywriter. It is a sequence designer that runs branching, behavior-triggered, multi-step paths to outcome. Deploy it after the Buyer Persona and Content Strategy agents are live.

- 01** Inventory every nurture sequence currently running. Audit them for branching logic, exit criteria, and outcome attribution. Most companies have nurture sequences that do not exit. The agent will not deploy into that mess until it is cleaned.
- 02** Define the sequence taxonomy. Post-content-download. Post-event. Post-demo. Win-back. Vertical-specific. Each type has a defined goal, defined exit criteria, and a defined success metric. The agent works within this taxonomy.
- 03** Pick three sequences for first deployment. Rebuild them with the agent. Compare against the existing versions in a holdout test for 60 days.
- 04** Promote winning rebuilds to live state. Sequences where the agent did not win get investigated, retrained, or retired.
- 05** Long-run on a quarterly sequence-refresh cadence. Email decay is real. Sequences that worked last quarter need retraining.

SUCCESS METRICS • WHAT TO MEASURE

Sequence-to-meeting rate	Percent of contacts entering a sequence that convert to a sales meeting within 90 days. Target: 2x pre-agent baseline.
Exit-criterion integrity	Percent of contacts who satisfy the exit criterion before the sequence ends without forced timeout. Target: > 60%.
Spam-complaint rate	Percent of recipients reporting any sequence step as spam. Target: < 0.1%. Higher than 0.3% is a kill criterion.
Sequence ROI	Pipeline attributable to nurture sequences per dollar of email program cost. Target: > 5x.

The Intent Signal Analysis Agent

PURPOSE. Synthesizes first-party and third-party intent signals into account-level "what changed and what to do" briefs. Distinguishes high-signal events from noise. [Most intent data is mostly hype](#); this agent's job is to enforce that judgment programmatically.

INPUTS. First-party intent (web visits, content engagement, demo requests, pricing-page hits); third-party intent feeds; news and funding feeds; LinkedIn engagement; community signals; product-usage signals where applicable.

OUTPUTS. Daily account-level signal briefs (top 20 accounts with new signals); signal-to-action recommendations (what motion fits each signal type); a weekly "signal noise" report flagging which third-party feeds added value vs. just volume.

MEMORY. Signal-to-outcome correlation history (which signals predicted pipeline, which did not); account-level signal history; signal-decay calibration (how quickly each signal type loses predictive value).

HUMAN OVERSIGHT. ABM lead reviews top-account signals daily. Quarterly review of signal-source ROI — sources that fail to predict pipeline are dropped.

COMMON MISTAKES. Treating volume of signals as a signal. Failing to calibrate signal decay. Buying a third-party intent feed and never auditing whether it added pipeline.

HOW TO DEPLOY

The Intent Signal Analysis Agent converts noisy intent feeds into prioritized action. Without this agent, intent data is a distraction. With it, intent data becomes the highest-leverage input to the Account Selection Agent and the ABM Orchestration Agent.

- 01** Define what counts as a meaningful intent signal for your business. Generic content topic intent rarely qualifies. Buyer-side comparison searches, RFP-language queries, and review-site visits qualify. The agent operates only on signals you have defined as meaningful.
- 02** Connect the agent to your first-party signals (gated content, demos, ad clicks), at least one third-party intent feed, and the Account Selection Agent's tier list. Intent signals on non-ICP accounts get filtered out, not investigated.
- 03** Run the agent in advisory mode for three weeks. It produces a weekly Top 20 signal report with recommended action per account. ABM lead reviews and decides what is acted on. The agent does not yet trigger workflows.
- 04** After three weeks, promote the agent to trigger ABM Orchestration motion on the highest-confidence Tier 1 signals. Lower-confidence signals stay in advisory mode for further development.
- 05** Tune the signal-to-action thresholds monthly. Every month the agent reports false-positive rate, missed-signal rate, and time-from-signal-to-pipeline. Tune the thresholds against this report.

SUCCESS METRICS • WHAT TO MEASURE

Signal-to-pipeline conversion

Percent of agent-flagged signals that produce a meeting or pipeline entry within 30 days. Target: > 12%.

False positive rate

Percent of high-confidence signals where ABM motion produced no engagement. Target: < 30%.

Time-to-action

Median time from signal detection to first outbound touch. Target: < 48 hours for Tier 1.

Coverage

Percent of pipeline-creating accounts where the agent caught a signal in the 60 days prior to opportunity creation. Target: > 50%.

The Hero *Workflow.*

HERO WORKFLOW

Tier 1 ABM motion design.

How an agent stack runs Tier 1 motion across the accounts that matter.

Eight agents and two human checkpoints turn a Tier 1 account list into a coordinated motion that executes across content, ads, outbound, sales, and intent monitoring. The ABM lead edits motion plans rather than building them. The agent system runs the orchestration. The human runs the strategy.

Eight agents, two human checkpoints, one orchestrated Tier 1 motion. The ABM lead's job becomes editing motion plans, not building them. The agent system runs the orchestration so the human can run the strategy.

AGENTS INVOLVED

- Account Selection Agent (maintains the Tier 1 list)
- Intent Signal Analysis Agent (triggers motion on Tier 1 signals)
- ABM Orchestration Agent (executes the per-account playbook)
- Competitor Intelligence Agent (informs account-specific messaging)
- Content Strategy Agent (provides asset library)
- Sales Enablement Agent (produces per-account materials)
- Email Nurture Sequence Agent (runs the long-tail track)
- Analytics & Reporting Agent (tracks engagement and pipeline)

HUMAN CHECKPOINTS

- *Joint weekly Tier 1 review with Sales VP (account additions, removals, promotions)*
- *Per-motion plan review before the motion goes live*

THE WORKFLOW, STEP BY STEP

- 01** List maintenance. The Account Selection Agent updates the Tier 1 list weekly off ICP grades, intent signals, and account-level changes (leadership moves, funding rounds, technology adoption). The ABM lead and Sales VP run a joint review every Monday.
- 02** Trigger detection. The Intent Signal Analysis Agent monitors high-confidence buying signals on Tier 1 accounts and notifies the ABM Orchestration Agent.
- 03** Motion plan generation. The ABM Orchestration Agent assembles the per-account motion plan: content sequence, ad targeting, sales touchpoint cadence, escalation triggers. The plan is account-specific, not template-rendered.
- 04** Asset production. The Sales Enablement Agent produces account-specific materials. The Content Strategy Agent flags any cornerstone or supporting content relevant to the account context. The Competitor Intelligence Agent provides current-state competitive positioning.
- 05** Motion execution. The ABM Orchestration Agent runs the motion across channels. Outbound, ads, content delivery, sales nudges. Every touch is logged to the account record.
- 06** Engagement tracking. The Analytics & Reporting Agent measures account-level engagement, scores intent escalation, and flags accounts that are heating up or going dark.
- 07** Sales handoff. When an account hits engagement thresholds, the named AE is notified with a current-state summary, brief, and recommended next action.

SUCCESS METRICS

- Tier 1 motion fidelity. Target: > 90% of planned touches execute on schedule.
- Time from signal to first touch. Target: < 48 hours.
- Tier 1 engagement rate. Target: > 25% of accounts produce a measurable engagement within 60 days.
- Tier 1 pipeline creation. Target: trend up, with a defined baseline by quarter 1.
- Sales co-ownership. Target: 100% of Tier 1 accounts have a named AE actively participating.

COMMON FAILURE MODES

- Marketing-side ABM that sales does not know about. Fix: the joint Monday review is non-negotiable. Without it, the program fails inside one quarter.
- Template motion across all Tier 1 accounts. Fix: the ABM Orchestration Agent must produce per-account plans, not templated runs. If every account gets the same touches, this is not ABM, it is mass marketing.
- Engagement signal hoarding. Fix: when intent escalates, sales is notified the same day. Marketing-side engagement reporting that never reaches the AE is a process failure.

Pre-flight *checklist.*

Run through this before activating a single agent in the demand stack. Each item is a precondition. If any item is unresolved, the agent that depends on it will produce noise, not pipeline.

- ☒ The ICP is defined across four dimensions (firmographic, behavioral, technographic, situational), not just firmographic. Without this, intent and campaign optimization optimize against the wrong target.

- ☒ Marketing-sourced pipeline is the agreed primary metric between marketing and sales leadership, with a written definition both sides have signed.

- ☒ Attribution model is documented and adopted. First-touch, multi-touch, or weighted, the model matters less than agreement on what model the team is running.

- ☒ Campaign budget authority is established. The Campaign Optimization Agent will need autonomy thresholds. Decide them before the agent goes live, not after the first surprise budget shift.

- ☒ Email program is in compliance health (DMARC, DKIM, SPF properly configured, unsubscribe path tested, suppression lists current). The Email Nurture Sequence Agent inherits whatever email reputation already exists.

- ☒ CRM lead-to-account matching works. Without it, demand pipeline never traces back to ICP-graded accounts and the agents work blind to the strategy.

- ☒ Sales is co-aligned on the Tier 1 ABM playbook the ABM Orchestration Agent will run. ABM orchestration without sales adoption is broadcast email at a higher cost.

The 30-60-90 *Roadmap.*

A demand-side rollout sequence for the first 90 days. The order matters. Each phase enables the next.

DAYS 1-30

Stand up the Foundation. ICP definition, attribution model, metrics dictionary, voice document. Do not deploy a single demand agent yet. Most demand teams want to skip this phase. They are the teams that fail. The Foundation is the substrate every demand agent reads from. A weak Foundation produces weak agent output, indefinitely. Audit the existing campaign infrastructure (paid platforms, marketing automation, attribution wiring) and document what works, what does not, and what is undocumented. Most undocumented integrations are the source of future agent failures. Resolve the pre-flight checklist above.

DAYS 31-60

Deploy in advisory mode. The Intent Signal Analysis Agent and the Campaign Optimization Agent both go live in advisory mode this phase. They produce daily recommendations to the demand lead. The demand lead manually approves or rejects every recommendation. The point is calibration, not autonomy. By day 60, the demand lead should have a felt sense for when the agents are right and when they are wrong. That felt sense becomes the trust foundation for bounded autonomy in the next phase. Start drafting the Tier 1 ABM playbook the ABM Orchestration Agent will run, in joint workshops with the named ABMer and the Sales VP.

DAYS 61-90

Promote to bounded autonomy. The Campaign Optimization Agent gets a defined daily-shift budget. The Intent Signal Analysis Agent triggers ABM Orchestration motion on high-confidence Tier 1 signals. The Email Nurture Sequence Agent ships three rebuilt sequences against held-out controls. Demand lead audits weekly, not daily. Begin live ABM Orchestration on a held-out sample of fifteen Tier 1 accounts.

DAY 91+

Expand and harden. Promote ABM Orchestration to full Tier 1 coverage once motion fidelity holds. Tune autonomy thresholds quarterly off the audit log. Begin the governance discipline of risk-tiered review, kill criteria, and quarterly retrospectives that this Field Note describes in the closing chapter.

Pitfalls & *remediation.*

PITFALL 01

Skipping the Foundation phase.

The most common failure pattern. Demand leaders see the agent layer and want to ship. They skip the ICP, the metrics, and the attribution work because it feels slow. Then every agent they deploy optimizes against the wrong target. The Foundation phase is non-negotiable.

PITFALL 02

Optimizing for clicks instead of pipeline.

The Campaign Optimization Agent will optimize for whatever metric you tell it to. Most teams default to platform-level metrics because they are easy to wire up. That is the trap. The agent must optimize for qualified pipeline, attributed back through the attribution model.

PITFALL 03

Treating ABM as a list, not a playbook.

ABM Orchestration without a defined playbook becomes mass email at a higher cost per touch. The playbook is what makes Tier 1 motion different from Tier 4 motion. The agent runs the playbook. If the playbook does not exist, the agent has nothing to run.

WHERE THIS FITS

Where this fits.

The four agents and one workflow in this playbook are the demand-generation slice of a larger marketing operating system. The full architecture has five layers, twenty-five agents, four hero workflows, and a governance model that earns trust at the executive level.

If you want the complete picture, read the flagship playbook: **AI Agents for Marketing Teams**. Free download. No signup. Send it to your team.

— Erik R. Miller

Erik R. Miller

SENIOR B2B MARKETING EXECUTIVE • FRACTIONAL CMO

Web: erikrmiller.com

Newsletter: The Operator, theoperator-erm.beehiiv.com

LinkedIn: linkedin.com/in/erikrm

Facebook: facebook.com/ERMAdivisory

Instagram: instagram.com/ermadvisory

Advisory: erikrmiller.com/#contact

The function with the visible *pipeline number* needs an agent stack, not another tool.

This Field Note is one volume in the Operator Playbook series. The full system has 25 agents, four hero workflows, a governance model, and a 90-day rollout.

Six other volumes cover Demand, Content, ABM, Operations, ICP, and Sales Enablement. The flagship pulls all of it together.

NEXT STEP

Get the *full playbook*.

FLAGSHIP	Read the full 65-page playbook
SERIES	erikrmiller.com/resources
NEWSLETTER	theoperator-erm.beehiiv.com
FACEBOOK	facebook.com/ERMA Advisory
INSTAGRAM	instagram.com/ermadvisory
ADVISE	erikrmiller.com/#contact